

A fast and accurate zebra finch syllable detector

Ben Pearre, L. Nathan Perkins, Jeffrey E. Markowitz, Timothy J. Gardner

January 8, 2016

Abstract

We present an accurate, versatile, and fast syllable detector that can control hardware at precisely timed moments during zebra finch song. Most moments during song can be isolated and detected with $> 95\%$ accuracy, easier syllables exceed 99.5% detection, and false positive rates can be $< 1\%$. The detector can run on a stock Mac Mini with a triggering delay around 1 millisecond and trigger jitter with $\sigma \approx 3$ milliseconds.

How to explain how to quantify this in the abstract?

1 Introduction

The adult zebra finch sings a song made up of 2-6 syllables, with longer songs taking on the order of a second. The song may be repeated hundreds of times per day, and is almost identical each time. This consistency presents a unique opportunity to study the neural basis of learning, audition, and control: we do not know of another model in which neural activity may be so easily correlated with motor output.

In order to take advantage of this consistency, it is useful to be able to detect selected moments during song, in order to align data or trigger other systems. To the best of our knowledge, current systems require expensive hardware, extensive hand-tuning, and careful choice of an easy syllable.

We developed a standalone detector based on the song spectrogram. Given a set of aligned songs, our software computes spectrograms and trains a neural network to output a TTL pulse at the chosen moment. The approach consists of three steps:

1. Align songs. This is outside the scope of this paper, but we provide a pointer to our song alignment software in Section A.
2. Choose one or more desired trigger syllables, and train a neural network to recognise them. This is carried out offline using any recent version of Matlab, and we recommend a computer with at least 32GB of RAM.
3. Once trained, the Matlab output file is used by a realtime detection program that listens to an audio signal and indicates detection of the target syllables via a TTL pulse. We present three implementations that trade off hardware requirements, ease of maintenance, and performance.

Quantify? Remove vacuous handwaving? Figure?

Examples?

Terminology: “syllable” is a unit, but this detector triggers on “song moments” or something?

We present the method in Section 2. Section 3 describes how we define and test performance. Section 4 presents our measurements. Section 5 gives an example usage case of the detector. We conclude in Section 6, and point to software resources in Appendix A.

2 Method

2.1 Learning a detector

We begin with a few hundred time-aligned recordings of a given bird’s song.

One or more moments during the song must be chosen. Our interface presents the time-aligned spectrogram averaged over training songs, and requires manual input of the target times. The user may next tune the region in frequency and time used for syllable identification. Then we assemble the training set from the song data, train the network, compute optimal output unit thresholds, and save the network file and an audio test file.

2.1.1 Recognition region

The neural network uses a contiguous set of the most recent timesteps (frames) from the song spectrogram. The frequencies F and timesteps T of this recognition region should be chosen in order to contain unique features of the target syllable and surrounding areas. This process could perhaps be automated, but we have not done so, as hand-choosing is neither difficult nor particularly error-prone.

2.1.2 Song micro-realignment

As conditions change, and especially during undirected song, syllable length and relative timing may vary slightly, which introduces variations in the precise timing of each syllable. The alignment software we use ensures that songs are aligned at the point midway through the song, but if the target syllable is not at that point, it is helpful to re-align the songs at the point of interest. This may be accomplished by looking for peaks in the correlation of the time-domain signal with the song whose spectrogram is closest to average over the training set.

2.1.3 Normalisation

To accommodate differences in amplitude, due to changes in the relative position between the microphone and bird or due to other minor variations in recording, normalisation ensures that the detector is sensitive only to the relative spectral content. We found that a two-step normalisation works. The spectrogram in the time and frequency windows, in decibels, is normalised to have mean zero and unit standard deviation. This eliminates relative differences in song amplitude due to microphone position or line levels. Next, each time-frequency bin is

I hope that’s the right thing to do. Seems to work, anyway...

What data is each norm computed over? Why does this do more than just the second step?

normalised based on the training set, such that that bin has mean zero and unit variance. This scales each time-frequency bin based on the amplitudes seen in the training set. These two normalisation steps provide a set of inputs that were more robust to outliers and less likely to produce false positives during silence when evaluated against other normalisation schemes, such as linear or L2 normalisation.

Would it be worth putting in data to this effect?

2.1.4 Building the training set

The neural network’s training set is created in the typical fashion: the rectangular $|F| \times |T|$ recognition region in the spectrogram is simply reshaped into a vector of length $|F||T|$. The frequency range is constant, and every possible range of the correct length for T is converted into a training input vector.

awkward phrasing

Training targets are, roughly, 1 if the input vector comes from the target time, 0 otherwise, for each target syllable. Since the song realignment may not be perfect, due to sample aliasing, and because the song spectrogram appears not to vary faster than the frame rate we chose, a strict binary target may ask the network to learn that practically identical samples should have opposite targets. Thus it is preferable to spread the output vector in time, such that at the target moment it is 1, and at neighbouring moments it is nonzero. We found that a Gaussian smoothing kernel around the target time with $\sigma \simeq 3\text{ms}$ serves well.

With inputs well outside the space on which a neural network has been trained, its outputs will be essentially random. In order to reduce the false positive rate it is useful to provide negative training examples that include silence, cage noise, non-song vocalisations, and perhaps songs from other birds. We have found that training with as low as a 1:1 ratio of non-song to song yields excellent results, although this will depend on the makeup of the non-song data.

2.1.5 Training the network

The network is trained using Matlab’s neural network toolbox. We tried a variety of feedforward neural network geometries, from simple 1-layer perceptrons to more complicated forms and many hidden nodes. Perhaps surprisingly, even the former yields excellent results on many syllables, but a 2-layer perceptron with a very small hidden layer—just 1 or 2 more units than the number of target syllables—was a good compromise between accuracy and training speed. Various other neural network geometries could be tried, as well as any other classifier that executes quickly. For more variable songs, deep structure-preserving networks may be more appropriate, but they are slow to train and unnecessary for zebra finch song.

Matlab’s neural network toolbox defaults to Levenburg-Marquardt training. This is a fast algorithm, but is memory-intensive, so multiple output syllables or high FFT frame rates require a large amount of RAM and increase training time to hours. Other training algorithms that use less RAM are much slower, and

Didn’t you guys try reducing input size with an autoencoder? Wasn’t that effective?

by default they will often terminate before converging due to their performance gradient going to 0.

2.1.6 Computing optimal output thresholds

When the network is trained, outputs of the classifier for any input are now available, and will be in the interval $(0, 1)$. We must choose a threshold above which the output is considered a positive detection. Finding the optimal threshold requires two choices. The first is the relative cost of false negatives to false positives, C . The second is the acceptable time interval: if the true event occurs at time t , and the detector triggers at any time $t \pm \Delta t$, then it is considered a correct detection. Then the optimal detection threshold τ is the one that minimises $[\text{false positives}] + C \cdot [\text{false negatives}]$ over the training set, using the definitions of false positives and negatives given in Section 3.1. Since large portions of the cost function are flat, random-restart hillclimbing would be effective, but a brute-force search requires fractions of a second. For the results presented here, we have used $C = 1$ and $\Delta t = 20\text{ms}$.

Or is it $(-1, 1)$?

2.1.7 Our parameter choices

We use a size-256 FFT, a Hamming window, and a spectrum sampled every 3 milliseconds. We usually define the network's input space to be 20-80ms long, and to span frequencies from about 1.5-7kHz, which contains the fundamentals and several overtones of most zebra finch vocalisations.

We found these parameters to work well across a variety of target syllables, but various other parameter sets yield results similar to those presented here. Some of the parameters trade off detection accuracy or detection timing vs. training time. For example, detection latency and jitter cannot be less than the FFT frame rate, but increasing the frame rate increases the size of the training set and thus the training resources required and the detection latency.

I chose this large value for Δt before I did target syllable realignment, and it could probably be much reduced now. Does this number make our results look bad? Should I explain it, change it and rerun, etc?

Check with Nathan and Jeff on what they've used.

The parameters used varied from iteration to iteration. Towards the end, we were using either 128 or 256 size FFT with no overlap (or even a gap between the FFTs). Jeff may be able to give a better sense of what worked best for him. - NP

2.2 Realtime detection

We tested three implementations of the realtime detector.

2.2.1 Matlab

This detector uses the built-in audio input and output jacks on a compatible computer. We tested on a 2015 Mac Mini. The 2015 Mac Pro does not have an audio input jack, so we tested using an M-Audio MobilePre external USB audio interface, but this did not improve measured latency, perhaps due to USB data transfer overhead.

The most versatile implementation, Matlab runs on a variety of hardware and operating systems, and is the easiest to modify. However, it performed relatively poorly on our test systems. Key to jitter performance is the size of the audio buffer, which on a 2015 Mac Pro or Mac Mini running Matlab 2015B

the minimum buffer size to prevent read underruns was about 4ms. Latency is also significantly longer than for the other two detectors.

2.2.2 Swift

This uses the Swift programming language and CoreAudio interface for Apple’s operating systems.

The Core Audio frameworks provide an adjustable buffer size for reading from and writing to audio hardware. Tuning this buffer size provides a tradeoff between the jitter in the detection and the processor usage needed to run the detector. We settled on a buffer size of 32 samples (0.7ms at 44.1kHz), as this created minimal system load while achieving detection within the desired lag and jitter.

Audio samples from the microphone are appended to a circular buffer. When there are sufficient samples in the buffer, the spectrogram is calculated by applying the Hamming window, performing an FFT and extracting the power. All vector operations are performed using the Accelerate framework (vDSP and vecLib), which use modern vector-oriented processor instructions to perform calculations. Outputs of the spectrogram, from the target frequency band, are appended to a second circular buffer. When there is sufficient data in the second circular buffer, the data is sent to a static implementation of the neural network previously trained, again implemented using the Accelerate framework to optimize normalization, matrix multiplication and all transfer functions.

The output of the neural network is written to the audio output buffer (again, in 32 sample or 0.7ms chunks) to create a trigger pulse, which can be used to trigger experimental feedback.

The above code can be run on multiple channels of audio on consumer hardware (such as a 2015 Mac Mini) with little impact on CPU usage ($< 15\%$). Depending on the experimental needs, latency can be further decreased (at the expense of processor usage) by adjusting the audio buffer sizes.

2.2.3 LabView

This implementation requires special software and hardware: LabView from National Instruments, and a data acquisition card—we use the National Instruments PCI-6251 card on a Windows PC.

This implementation is difficult to modify and debug, and requires Windows. However, our test configuration achieved excellent performance, and further gains should be possible if the implementation were retargeted onto FPGA hardware.

3 Quantification

Ground truth is given on the training set, and can be measured by presenting the recorded training songs as well as the canonical detection events. To this end, besides the trained network object, our learning code produces an audio

file consisting of all of the training data on the left audio channel and a delta function at each moment of correct detection (basically a TTL pulse, although the voltage will fluctuate with audio output volume and the exact shape of the pulse will vary with the quality of the audio player) on the right channel. Thus, when played on any audio player, the left channel may be provided as input to the detector, and the right channel may be compared against the detector’s detection pulse.

3.1 Accuracy

The Matlab neural network toolbox breaks the given training set into three groups: data on which the network is trained, data used to validate the progress of the training algorithm, and holdout test data used only as a final measure of performance. We further withhold a portion of the training data in order to provide another evaluation of performance on unseen data.

We define the accuracy of the network based on its classification performance per frame. In order to avoid the apparent problem of every non-detected non-syllable counting as a true negative, we also tried defining accuracy on a per-song basis, such that a song without the target syllable counted as a single true negative. Computing the optimal output thresholds on a per-frame basis resulted in higher thresholds and thus a lower false-positive rate, with minimal consequences to the false-negative rate.

The accuracy as defined above is used for computing the optimal thresholds above which the network’s output should be interpreted as a match on the training data as described in Section 2.1.6, for evaluation of the detector on the training songs in Matlab and Swift, and while live.

We present the resulting confusion matrix for a few sample songs, and, for simplicity, we present a summary of detector accuracy using the area under the ROC curve.

3.2 Timing

We evaluate the time taken from the presentation of the target syllable to the firing of the detector’s TTL pulse. While playing the audio test file from another device (such as a mobile phone), the TTL output from the ground-truth channel of the audio output may be used as the trigger pulse for an oscilloscope, and compared to the TTL pulse produced by the Swift detector, which sees only the birdsong channel of the audio file. For this purpose we used a pulse generator (Philips PM 5715) to widen the detector’s output spike to a number much larger than the jitter (100ms). This obviates pulse length jitter in the output device by essentially discarding the falling edge of the output pulse. The oscilloscope is then set to averaging mode (128-trigger average) in order to collect timing data. The canonical signal is the trigger at $t = 0$, and the average of the detector’s detection events will be seen as a low-to-high transition with form approximating the cumulative probability distribution function (CDF) of the detector’s output in response to the chosen song event.

It’s a little more complex than that (cross-validation), but can we go with this?

But another test set layer is just paranoia!

Listed latency of 50ns, jitter of $\leq 0.1\%$ or 50ps, whichever is greater.




Figure 1: A spectrogram made by averaging over 526 songs. The 6 target syllables, spaced 10ms apart, are marked by red lines.

Mean latency is then given as the halfway point of that detection sigmoid. It is a helpful number, but not a critical one, since a detector with high but constant latency can be trained to trigger at a point somewhat before the true moment of interest. Often more important is latency jitter: how much random variability is there in the latency?

We obtain both of these numbers by performing a maximum-likelihood fit of a Gaussian distribution to the timing data obtained from the oscilloscope.

4 Results

Figure 1 shows a typical song, with six target moments selected. We trained the network with 8 hidden units for the 6 detection targets spaced 10ms apart. Matlab’s self-report of detection performance for each iteration of the song is shown in Figure 2, with ROC curves shown in Figure 3. The beginning of the syllable is difficult to detect, with a few detection events considerably earlier than the correct moment. But the syllable quickly becomes reliably identifiable. By the time the detector has seen 50ms of the syllable—the sixth detection point—performance is good: the area under the ROC curve is 0.992.

Figure 4 shows the latency and jitter for both the LabView implementation and the Swift implementation detector running on a single syllable from another bird, using the test audio file generated during training on that song (not shown). The latency and jitter were measured using an oscilloscope triggering on the true signal, using a pulse generator as described above (a single detection event, non-averaged, would be a single low-to-high step function). For the Swift implementation, the average detection latency is 6.5ms, while latency jitter has $\sigma = 1.7\text{ms}$ (MLE estimate). For the LabView implementation, latency is 4.6ms, while latency jitter has $\sigma = 1.4\text{ms}$.

Figure ?? could show a comparison between Jeff’s filtering approach on the TDT and my FFT-NN. But how much time do I want to spend describing Jeff’s filters? Not much, since it’s not in use anywhere besides our lab. Or I could simply insert a photo of Jeff, holding my FFT-NN and a finch, and smiling, with “My name is Dr. Jeffrey E. Markowitz and I endorse this software!”

Bases for comparison?
Jeff’s matched filterbank?
Nothing really
informative.
I should generate, say, 100
random detection points
in a few songs, so I can
say what something like
“average accuracy” is.

FIXME

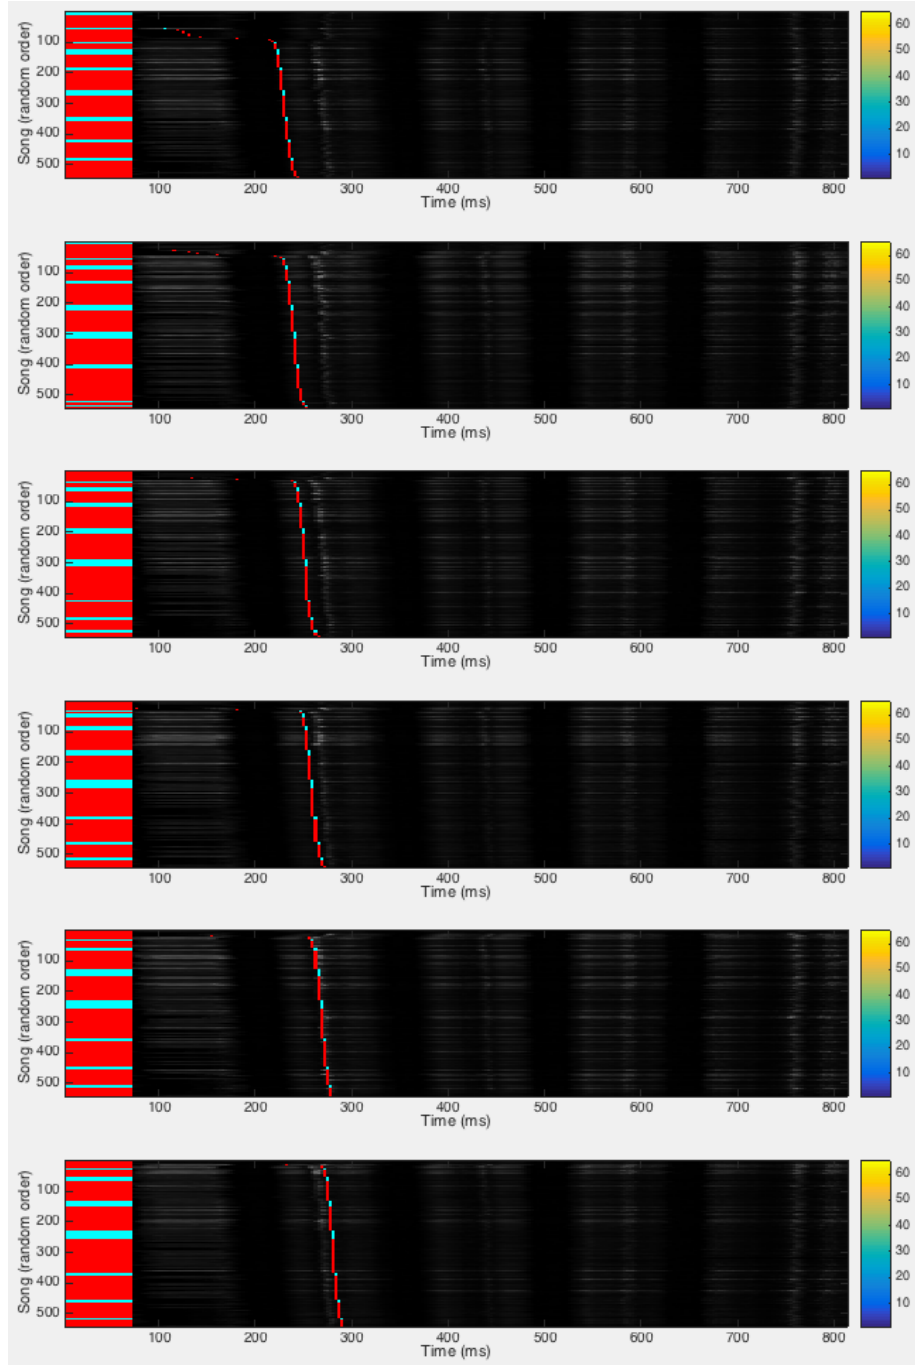


Figure 2: Detection results. Each plot shows detection events for the corresponding syllable as shown in Figure 1. The horizontal axis is time, and the vertical axis is the index of the 526 single song presentations. The grey shading shows the total audio energy of song Y at time X. The bar on the left is the same width as the detection window, so no detection events can happen within that region, and its colour code shows training songs to the right of red regions and unseen test songs to the right of cyan regions. For visualisation of the distribution, songs have been stably sorted by the time of detection events.




Figure 3: ROC curves for the detection of targets shown in Figure 1. The first one, at the beginning of a syllable, is the most difficult to detect, and as more structure emerges in the current syllable, accuracy approaches 100%.

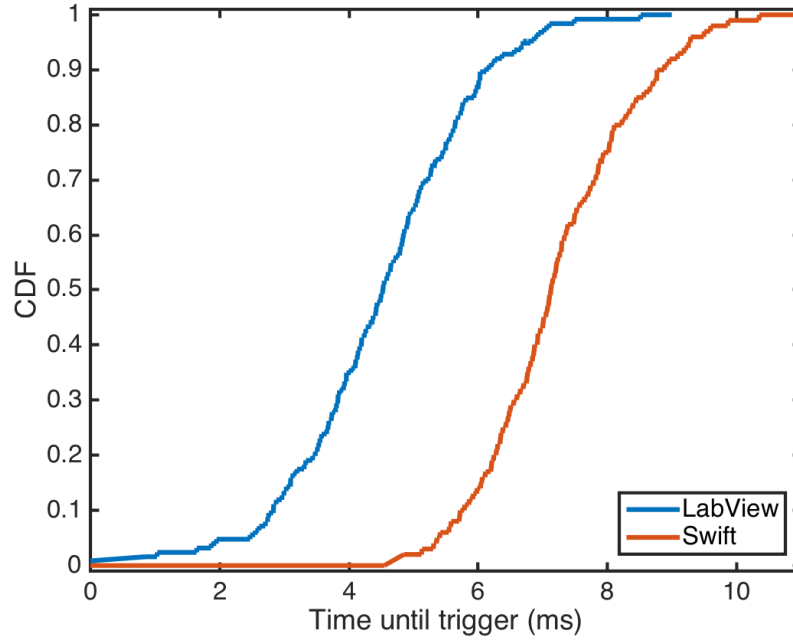


Figure 4: Timing curve for a syllable running a test file generated by training.

5 Case Study?

Jeff is using the detector for something cool, and it might be fun to get him to describe it in a paragraph or two.

Failing that, I can immediately do song-aligned X spike rasters if lw95rhp is still willing to sing... but I have not heard him do so lately.

Do we want to include a case study? I think the paper is on the long side already, but it might be interesting and let Jeff earn that coauthorship ;)

6 Conclusion

This syllable detector is appropriate for zebra finch song. It offers the following benefits:

- False positive and false negative rates can be well under 1%.
- Latency and jitter are both in the range of 2 milliseconds.
- Works on a wide range of target syllables.
- Requires minimal hand-tuning.
- Runs in realtime on inexpensive consumer-grade hardware—we use the Mac Mini with 8GB RAM.
- We recommend training the network on a computer with at least 32GB of RAM.

It has not been evaluated in other species, and the high accuracy presented here relies on zebra-finch-like consistency of song.

A Resources

Song alignment: Last we checked, Jeff Markowitz’s song alignment software could be found at
<https://github.com/jmarkow>.

Training the neural network: Our implementation of the syllable detector training code is available under the GNU GPL at:
<https://github.com/bwpearre/>

Fix the location! If going live, I should reorganise my repository, and also clean up the code.

Runtime: The Swift implementation for executing the trained network:
<https://github.com/nathanntg/syllable-detector>

References